

```
#Performing Twitter Hate Speech Recognition using LSTM(Long Short Term Model)
```

```
import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
import nltk
from nltk.corpus import stopwords
import re
```

```
#Loading the Sentiment 140 dataset
```

```
data =
pd.read_csv('/Users/Andrea/Desktop/training.1600000.processed.noemoticon.csv', encoding='latin-1', header=None)
data.columns = ['target', 'ids', 'date', 'flag', 'user', 'text']

data.head(50)
```

	target	ids	date	flag	\
0	0	1467810369	Mon Apr 06 22:19:45 PDT 2009	NO_QUERY	
1	0	1467810672	Mon Apr 06 22:19:49 PDT 2009	NO_QUERY	
2	0	1467810917	Mon Apr 06 22:19:53 PDT 2009	NO_QUERY	
3	0	1467811184	Mon Apr 06 22:19:57 PDT 2009	NO_QUERY	
4	0	1467811193	Mon Apr 06 22:19:57 PDT 2009	NO_QUERY	
5	0	1467811372	Mon Apr 06 22:20:00 PDT 2009	NO_QUERY	
6	0	1467811592	Mon Apr 06 22:20:03 PDT 2009	NO_QUERY	
7	0	1467811594	Mon Apr 06 22:20:03 PDT 2009	NO_QUERY	
8	0	1467811795	Mon Apr 06 22:20:05 PDT 2009	NO_QUERY	
9	0	1467812025	Mon Apr 06 22:20:09 PDT 2009	NO_QUERY	
10	0	1467812416	Mon Apr 06 22:20:16 PDT 2009	NO_QUERY	
11	0	1467812579	Mon Apr 06 22:20:17 PDT 2009	NO_QUERY	
12	0	1467812723	Mon Apr 06 22:20:19 PDT 2009	NO_QUERY	
13	0	1467812771	Mon Apr 06 22:20:19 PDT 2009	NO_QUERY	
14	0	1467812784	Mon Apr 06 22:20:20 PDT 2009	NO_QUERY	
15	0	1467812799	Mon Apr 06 22:20:20 PDT 2009	NO_QUERY	
16	0	1467812964	Mon Apr 06 22:20:22 PDT 2009	NO_QUERY	
17	0	1467813137	Mon Apr 06 22:20:25 PDT 2009	NO_QUERY	
18	0	1467813579	Mon Apr 06 22:20:31 PDT 2009	NO_QUERY	
19	0	1467813782	Mon Apr 06 22:20:34 PDT 2009	NO_QUERY	
20	0	1467813985	Mon Apr 06 22:20:37 PDT 2009	NO_QUERY	
21	0	1467813992	Mon Apr 06 22:20:38 PDT 2009	NO_QUERY	
22	0	1467814119	Mon Apr 06 22:20:40 PDT 2009	NO_QUERY	
23	0	1467814180	Mon Apr 06 22:20:40 PDT 2009	NO_QUERY	
24	0	1467814192	Mon Apr 06 22:20:41 PDT 2009	NO_QUERY	

25	0	1467814438	Mon Apr 06 22:20:44 PDT 2009	NO_QUERY
26	0	1467814783	Mon Apr 06 22:20:50 PDT 2009	NO_QUERY
27	0	1467814883	Mon Apr 06 22:20:52 PDT 2009	NO_QUERY
28	0	1467815199	Mon Apr 06 22:20:56 PDT 2009	NO_QUERY
29	0	1467815753	Mon Apr 06 22:21:04 PDT 2009	NO_QUERY
30	0	1467815923	Mon Apr 06 22:21:07 PDT 2009	NO_QUERY
31	0	1467815924	Mon Apr 06 22:21:07 PDT 2009	NO_QUERY
32	0	1467815988	Mon Apr 06 22:21:09 PDT 2009	NO_QUERY
33	0	1467816149	Mon Apr 06 22:21:11 PDT 2009	NO_QUERY
34	0	1467816665	Mon Apr 06 22:21:21 PDT 2009	NO_QUERY
35	0	1467816749	Mon Apr 06 22:21:20 PDT 2009	NO_QUERY
36	0	1467817225	Mon Apr 06 22:21:27 PDT 2009	NO_QUERY
37	0	1467817374	Mon Apr 06 22:21:30 PDT 2009	NO_QUERY
38	0	1467817502	Mon Apr 06 22:21:32 PDT 2009	NO_QUERY
39	0	1467818007	Mon Apr 06 22:21:39 PDT 2009	NO_QUERY
40	0	1467818020	Mon Apr 06 22:21:39 PDT 2009	NO_QUERY
41	0	1467818481	Mon Apr 06 22:21:46 PDT 2009	NO_QUERY
42	0	1467818603	Mon Apr 06 22:21:49 PDT 2009	NO_QUERY
43	0	1467818900	Mon Apr 06 22:21:53 PDT 2009	NO_QUERY
44	0	1467819022	Mon Apr 06 22:21:56 PDT 2009	NO_QUERY
45	0	1467819650	Mon Apr 06 22:22:05 PDT 2009	NO_QUERY
46	0	1467819712	Mon Apr 06 22:22:06 PDT 2009	NO_QUERY
47	0	1467819812	Mon Apr 06 22:22:07 PDT 2009	NO_QUERY
48	0	1467820206	Mon Apr 06 22:22:13 PDT 2009	NO_QUERY
49	0	1467820835	Mon Apr 06 22:22:25 PDT 2009	NO_QUERY

	user	text
0	_TheSpecialOne_	@switchfoot http://twitpic.com/2ylzl - Awww, t...
1	scotthamilton	is upset that he can't update his Facebook by ...
2	mattycus	@Kenichan I dived many times for the ball. Man...
3	ElleCTF	my whole body feels itchy and like its on fire
4	Karoli	@nationwideclass no, it's not behaving at all....
5	joy_wolf	@Kwesidei not the whole crew
6	mybirch	Need a hug
7	coZZ	@LOLTrish hey long time no see! Yes.. Rains a...
8	2Hood4Hollywood	@Tatiana_K nope they didn't have it
9	mimismo	@twittera que me muera ?
10	erinx3leannexo	spring break in plain city... it's snowing

11	pardonlauren	I just re-pierced my ears
12	TLeC	@caregiving I couldn't bear to watch it. And ...
13	robobbierobert	@octolinz16 It it counts, idk why I did either...
14	bayofwolves	@smarrison i would've been the first, but i di...
15	HairByJess	@iamjazzyfizzle I wish I got to watch it with ...
16	lovesongwriter	Hollis' death scene will hurt me severely to w...
17	armotley	about to file taxes
18	starkissed	@LettyA ahh ive always wanted to see rent lov...
19	gi_gi_bee	@FakerPattyPattz Oh dear. Were you drinking ou...
20	quanvu	@alydesigns i was out most of the day so didn'...
21	swinspeedx	one of my friend called me, and asked to meet ...
22	cooliodoc	@angry_barista I baked you a cake but I ated it
23	viJILLante	this week is not going as i had hoped
24	Ljelli3166	blagh class at 8 tomorrow
25	ChicagoCubbie	I hate when I have to call and wake people up
26	KatieAngell	Just going to cry myself to sleep after watchi...
27	gagoo	im sad now Miss.Lilly
28	abel209	ooooh.... LOL that leslie.... and ok I won't ...
29	BaptisteTheFool	Meh... Almost Lover is the exception... this t...
30	fatkat309	some1 hacked my account on aim now i have to ...
31	EmCDL	@alielayus I want to go to promote GEAR AND GR...
32	merisssa	thought sleeping in was an option tomorrow but...
33	Pbearfox	@julieebaby awe i love you too!!!! 1 am here ...
34	jsoo	@HumpNinja I cry my asian eyes to sleep at night
35	scarletletterm	ok I'm sick and spent an hour sitting in the s...
36	crosland 12	@cocomix04 ill tell ya the story later not a ...

```

37         ajaxpro  @MissXu sorry! bed time came here (GMT+1)  ht...
38         Tmttq86  @fleurylis I don't either. Its depressing. I d...
39  Anthony_Nguyen  Bed. Class 8-12. Work 12-3. Gym 3-5 or 6. Then...
40         itsanimesh  really don't feel like getting up today... but...
41         lionslamb  He's the reason for the teardrops on my guitar...
42         kennypham  Sad, sad, sad. I don't know why but I hate thi...
43  DdubsShellBell  @JonathanRKnight Awww I soo wish I was there t...
44         hpfangirl94  Falling asleep. Just heard about that Tracy gi...
45         antzpantz  @Viennah Yay! I'm happy for you with your job!...
46         labrt2004  Just checked my user timeline on my blackberry...
47         IrisJumbe  Oh man...was ironing @jeancjumbe's fave top to...
48         peacoats  is strangely sad about LiLo and SamRo breaking...
49         cyantist  @tea oh! i'm so sorry i didn't think about th...

```

```
data.tail(150)
```

	target	ids		date	flag	\
1599850	4	2193552448	Tue Jun 16 08:36:46 PDT 2009	NO_QUERY		
1599851	4	2193552454	Tue Jun 16 08:36:46 PDT 2009	NO_QUERY		
1599852	4	2193552500	Tue Jun 16 08:36:47 PDT 2009	NO_QUERY		
1599853	4	2193552501	Tue Jun 16 08:36:47 PDT 2009	NO_QUERY		
1599854	4	2193552510	Tue Jun 16 08:36:47 PDT 2009	NO_QUERY		
...	
1599995	4	2193601966	Tue Jun 16 08:40:49 PDT 2009	NO_QUERY		
1599996	4	2193601969	Tue Jun 16 08:40:49 PDT 2009	NO_QUERY		
1599997	4	2193601991	Tue Jun 16 08:40:49 PDT 2009	NO_QUERY		
1599998	4	2193602064	Tue Jun 16 08:40:49 PDT 2009	NO_QUERY		
1599999	4	2193602129	Tue Jun 16 08:40:50 PDT 2009	NO_QUERY		

	user	
text		
1599850	jenniehager	sick, sick, sick today, but still fingers cros...
1599851	YessicaCarcamo	Getting ready to go out I have to go to mac ...
1599852	afarley08	loves visitors so if you are goin to the Aeros...
1599853	Jay_Cousins	Nice out, but If I am going to leave here at

```

6...
1599854    laughingstars    @WillaRyan At least trees are being
spared.
...
...
1599995    AmandaMarie1028 Just woke up. Having no school is the best
fee...
1599996    TheWDBboards    TheWDB.com - Very cool to hear old Walt
interv...
1599997    bpbabe    Are you ready for your MoJo Makeover? Ask me
f...
1599998    tinydiamondz    Happy 38th Birthday to my boo of alll
time!!! ...
1599999    RyanTrevMorris    happy #charitytuesday @theNSPCC
@SparksCharity...

```

```
[150 rows x 6 columns]
```

```
data.shape
```

```
(1600000, 6)
```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1600000 entries, 0 to 1599999
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  ---
 0   target      1600000 non-null  int64
 1   ids         1600000 non-null  int64
 2   date        1600000 non-null  object
 3   flag        1600000 non-null  object
 4   user        1600000 non-null  object
 5   text        1600000 non-null  object
dtypes: int64(2), object(4)
memory usage: 73.2+ MB

```

```
data = data[['target', 'text']]
```

```
#Target values mapped to binary (0 = negative, 1 = positive)
```

```
#Ignore the neutral tweets (target = 2)
```

```
data = data[data['target'] != 2]
```

```
data['target'] = data['target'].map({0: 0, 4: 1})
```

```
stop_words = set(stopwords.words('english'))
```

```

def preprocess_text(text):
    text = re.sub(r'http\S+', '', text)    #Remove URLs
    text = re.sub(r'@\w+', '', text)       #Remove mentions

```

```

    text = re.sub(r'#', '', text)           #Remove hashtags
    text = re.sub(r'\d+', '', text)         #Remove numbers
    text = re.sub(r'[\W\s]', '', text)      #Remove punctuation
    text = text.lower()                     #Lowercase
    text = ' '.join([word for word in text.split() if word not in
stop_words])
    return text

data['cleaned_text'] = data['text'].apply(preprocess_text)

#Tokenize the text
tokenizer = Tokenizer(num_words=10000, oov_token='<OOV>')
tokenizer.fit_on_texts(data['cleaned_text'])
sequences = tokenizer.texts_to_sequences(data['cleaned_text'])

#Pad the sequences
padded_sequences = pad_sequences(sequences, padding='post', maxlen=50)

#Spitting the data
X_train, X_test, y_train, y_test = train_test_split(padded_sequences,
data['target'], test_size=0.2, random_state=42)

#Building the model
model = Sequential()
model.add(Embedding(input_dim=10000, output_dim=64, input_length=50))
model.add(LSTM(64, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(64))
model.add(Dropout(0.2))
model.add(Dense(1, activation='sigmoid'))

#Compiling the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

#Train the model
history = model.fit(X_train, y_train, epochs=5,
validation_data=(X_test, y_test), batch_size=32)

Epoch 1/5

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(

40000/40000 ————— 726s 18ms/step - accuracy: 0.4997 -
loss: 0.6933 - val_accuracy: 0.4984 - val_loss: 0.6933

```

```
Epoch 2/5
40000/40000 _____ 707s 18ms/step - accuracy: 0.4998 -
loss: 0.6932 - val_accuracy: 0.5016 - val_loss: 0.6932
Epoch 3/5
40000/40000 _____ 727s 18ms/step - accuracy: 0.5000 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932
Epoch 4/5
40000/40000 _____ 732s 18ms/step - accuracy: 0.4998 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932
Epoch 5/5
40000/40000 _____ 3283s 82ms/step - accuracy: 0.4995 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932
```

```
model.summary()
```

```
Model: "sequential"
```

Layer (type) Param #	Output Shape	
embedding (Embedding) 640,000	(32, 50, 64)	
lstm (LSTM) 33,024	(32, 50, 64)	
dropout (Dropout) 0	(32, 50, 64)	
lstm_1 (LSTM) 33,024	(32, 64)	
dropout_1 (Dropout) 0	(32, 64)	
dense (Dense) 65	(32, 1)	

```
Total params: 2,118,341 (8.08 MB)
```

Trainable params: 706,113 (2.69 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 1,412,228 (5.39 MB)

#Evaluating the model

```
predictions = (model.predict(X_test) > 0.5).astype("int32")
print(classification_report(y_test, predictions))
print('Accuracy:', accuracy_score(y_test, predictions))
```

```
10000/10000 ————— 5366s 537ms/step
              precision    recall  f1-score   support

     0         0.50         1.00         0.67     159494
     1         0.00         0.00         0.00     160506

 accuracy                   0.50     320000
  macro avg              0.25         0.50         0.33     320000
 weighted avg            0.25         0.50         0.33     320000
```

Accuracy: 0.49841875

```
/opt/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/opt/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/opt/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

#Performing Hate Speech Recognition in tweets using GRU

```
import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, GRU, Dense, Dropout
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
import nltk
```



```

from nltk.corpus import stopwords
import re

#Loading the Sentiment140 dataset
data =
pd.read_csv('/Users/Andrea/Desktop/training.1600000.processed.noemotic
on.csv', encoding='latin-1', header=None)
data.columns = ['target', 'ids', 'date', 'flag', 'user', 'text']

data = data[['target', 'text']]

#Target values mapped to binary (0 = negative, 1 = positive)
#Ignore the neutral tweets (target = 2)
data = data[data['target'] != 2]
data['target'] = data['target'].map({0: 0, 4: 1})

stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    text = re.sub(r'http\S+', '', text) #Remove URLs
    text = re.sub(r'@\w+', '', text) #Remove mentions
    text = re.sub(r'#', '', text) #Remove hashtags
    text = re.sub(r'\d+', '', text) #Remove numbers
    text = re.sub(r'^\w\s', '', text) #Remove punctuation
    text = text.lower() #Lowercase
    text = ' '.join([word for word in text.split() if word not in
stop_words])
    return text

data['cleaned_text'] = data['text'].apply(preprocess_text)

#Tokenize the text
tokenizer = Tokenizer(num_words=10000, oov_token='<OOV>')
tokenizer.fit_on_texts(data['cleaned_text'])
sequences = tokenizer.texts_to_sequences(data['cleaned_text'])

#Pad the sequences
padded_sequences = pad_sequences(sequences, padding='post', maxlen=50)

#Splitting the data
X_train, X_test, y_train, y_test = train_test_split(padded_sequences,
data['target'], test_size=0.2, random_state=42)

#Building the model
model = Sequential()
model.add(Embedding(input_dim=10000, output_dim=64, input_length=50))
model.add(GRU(64, return_sequences=True))

```

```

model.add(Dropout(0.2))
model.add(GRU(64))
model.add(Dropout(0.2))
model.add(Dense(1, activation='sigmoid'))

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(

```

#Compile the model

```

model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

```

#Training the model

```

history_1 = model.fit(X_train, y_train, epochs=5,
validation_data=(X_test, y_test), batch_size=42)

```

```

Epoch 1/5
30477/30477 ━━━━━━━━━━━ 696s 23ms/step - accuracy: 0.5003 -
loss: 0.6934 - val_accuracy: 0.4984 - val_loss: 0.6932
Epoch 2/5
30477/30477 ━━━━━━━━━━━ 704s 23ms/step - accuracy: 0.5009 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932
Epoch 3/5
30477/30477 ━━━━━━━━━━━ 708s 23ms/step - accuracy: 0.5102 -
loss: 0.6889 - val_accuracy: 0.7768 - val_loss: 0.4692
Epoch 4/5
30477/30477 ━━━━━━━━━━━ 714s 23ms/step - accuracy: 0.7821 -
loss: 0.4604 - val_accuracy: 0.7917 - val_loss: 0.4411
Epoch 5/5
30477/30477 ━━━━━━━━━━━ 716s 23ms/step - accuracy: 0.7987 -
loss: 0.4312 - val_accuracy: 0.7956 - val_loss: 0.4351

```

```

model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	
Param #		
embedding_1 (Embedding)	(None, 50, 64)	
640,000		
gru (GRU)	(None, 50, 64)	
24,960		

0	dropout_2 (Dropout)	(None, 50, 64)
24,960	gru_1 (GRU)	(None, 64)
0	dropout_3 (Dropout)	(None, 64)
65	dense_1 (Dense)	(None, 1)

Total params: 2,069,957 (7.90 MB)

Trainable params: 689,985 (2.63 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 1,379,972 (5.26 MB)

#Evaluating the model

```
predictions_1 = (model.predict(X_test) > 0.5).astype("int32")
print(classification_report(y_test, predictions_1))
print('Accuracy:', accuracy_score(y_test, predictions_1))
```

```
10000/10000 ————— 41s 4ms/step
              precision    recall  f1-score   support

     0       0.81         0.78         0.79      159494
     1       0.79         0.81         0.80      160506

 accuracy                   0.80         320000
  macro avg                 0.80         0.80         0.80      320000
 weighted avg               0.80         0.80         0.80      320000
```

Accuracy: 0.7956375

#Implementation of Bi-LSTM, Bi-GRU, LSTM Layer for classification of Sentiment 140 dataset for analysis

```
from tensorflow.keras.layers import BatchNormalization, Bidirectional
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout
from tensorflow.keras.layers import Embedding, GRU, Dense, Dropout
from sklearn.model_selection import train_test_split
```

```

from sklearn.metrics import classification_report, accuracy_score
import nltk
from nltk.corpus import stopwords
import re

#Build the model
model = Sequential()

#Embedding layer
model.add(Embedding(input_dim=10000, output_dim=128, input_length=50))

#First Bidirectional LSTM layer
model.add(Bidirectional(LSTM(64, return_sequences=True)))
model.add(Dropout(0.3))
model.add(BatchNormalization())

#Second Bidirectional GRU layer
model.add(Bidirectional(GRU(64, return_sequences=True)))
model.add(Dropout(0.3))
model.add(BatchNormalization())

#Third LSTM layer
model.add(LSTM(64))
model.add(Dropout(0.3))

#Dense layer for binary classification
model.add(Dense(1, activation='sigmoid'))

#Compiling the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(

#Train the model
history_2 = model.fit(X_train, y_train, epochs=5,
validation_data=(X_test, y_test), batch_size=32)

model.summary()

#Evaluate the model
predictions_3 = (model.predict(X_test) > 0.5).astype("int32")
print(classification_report(y_test, predictions_3))
print('Accuracy:', accuracy_score(y_test, predictions_3))

Epoch 1/5
40000/40000 ————— 5752s 144ms/step - accuracy: 0.5004 -

```

```

loss: 0.6937 - val_accuracy: 0.5016 - val_loss: 0.6931
Epoch 2/5
40000/40000 _____ 2809s 70ms/step - accuracy: 0.5002 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932
Epoch 3/5
40000/40000 _____ 2512s 63ms/step - accuracy: 0.4992 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932
Epoch 4/5
40000/40000 _____ 2496s 62ms/step - accuracy: 0.5007 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932
Epoch 5/5
40000/40000 _____ 2492s 62ms/step - accuracy: 0.4994 -
loss: 0.6932 - val_accuracy: 0.4984 - val_loss: 0.6932

```

Model: "sequential_2"

Layer (type) Param #	Output Shape	
embedding_2 (Embedding) 1,280,000	(32, 50, 128)	
bidirectional (Bidirectional) 98,816	(32, 50, 128)	
dropout_4 (Dropout) 0	(32, 50, 128)	
batch_normalization (BatchNormalization) 512	(32, 50, 128)	
bidirectional_1 (Bidirectional) 74,496	(32, 50, 128)	
dropout_5 (Dropout) 0	(32, 50, 128)	
batch_normalization_1 512	(32, 50, 128)	

(BatchNormalization)		
lstm_3 (LSTM)	(32, 64)	
49,408		
dropout_6 (Dropout)	(32, 64)	
0		
dense_2 (Dense)	(32, 1)	
65		

Total params: 4,510,405 (17.21 MB)

Trainable params: 1,503,297 (5.73 MB)

Non-trainable params: 512 (2.00 KB)

Optimizer params: 3,006,596 (11.47 MB)

10000/10000	155s 15ms/step			
	precision	recall	f1-score	support
0	0.50	1.00	0.67	159494
1	0.00	0.00	0.00	160506
accuracy			0.50	320000
macro avg	0.25	0.50	0.33	320000
weighted avg	0.25	0.50	0.33	320000

Accuracy: 0.49841875

```
/opt/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/opt/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/opt/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use
```

```

`zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

#Implementation of Bi-LSTM , Bi-GRU and Bi-Lstm layer for the
classification of tweets in Sentiment 140 dataset

from tensorflow.keras.layers import BatchNormalization, Bidirectional

#Build the model
model = Sequential()

#Embedding layer
model.add(Embedding(input_dim=10000, output_dim=128, input_length=50))

#First Bidirectional LSTM layer
model.add(Bidirectional(LSTM(64, return_sequences=True)))
model.add(Dropout(0.3))
model.add(BatchNormalization())

#Second Bidirectional GRU layer
model.add(Bidirectional(GRU(64, return_sequences=True)))
model.add(Dropout(0.3))
model.add(BatchNormalization())

#Third Bidirectional LSTM layer
model.add(Bidirectional(LSTM(64)))
model.add(Dropout(0.3))

#Dense layer for binary classification, activation function is sigmoid
model.add(Dense(1, activation='sigmoid'))

#Compiling the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(

#Training the model
history_3 = model.fit(X_train, y_train, epochs=5,
validation_data=(X_test, y_test), batch_size=32)

model.summary()

#Evaluate the model
predictions_4 = (model.predict(X_test) > 0.5).astype("int32")
print(classification_report(y_test, predictions_4))
print('Accuracy:', accuracy_score(y_test, predictions_4))

```

Epoch 1/5
 40000/40000 ————— 2899s 72ms/step - accuracy: 0.7664 -
 loss: 0.4842 - val_accuracy: 0.7906 - val_loss: 0.4463

Epoch 2/5
 40000/40000 ————— 2904s 73ms/step - accuracy: 0.7963 -
 loss: 0.4370 - val_accuracy: 0.7944 - val_loss: 0.4406

Epoch 3/5
 40000/40000 ————— 3843s 96ms/step - accuracy: 0.8035 -
 loss: 0.4237 - val_accuracy: 0.7955 - val_loss: 0.4374

Epoch 4/5
 40000/40000 ————— 4581s 115ms/step - accuracy: 0.8069 -
 loss: 0.4174 - val_accuracy: 0.7963 - val_loss: 0.4379

Epoch 5/5
 40000/40000 ————— 2950s 74ms/step - accuracy: 0.8097 -
 loss: 0.4132 - val_accuracy: 0.7944 - val_loss: 0.4390

Model: "sequential_3"

Layer (type) Param #	Output Shape	
embedding_3 (Embedding) 1,280,000	(32, 50, 128)	
bidirectional_2 (Bidirectional) 98,816	(32, 50, 128)	
dropout_7 (Dropout) 0	(32, 50, 128)	
batch_normalization_2 512 (BatchNormalization)	(32, 50, 128)	
bidirectional_3 (Bidirectional) 74,496	(32, 50, 128)	
dropout_8 (Dropout) 0	(32, 50, 128)	

batch_normalization_3	(32, 50, 128)	
512 (BatchNormalization)		
<hr/>		
bidirectional_4 (Bidirectional)	(32, 128)	
98,816		
<hr/>		
dropout_9 (Dropout)	(32, 128)	
0		
<hr/>		
dense_3 (Dense)	(32, 1)	
129		
<hr/>		

Total params: 4,658,821 (17.77 MB)

Trainable params: 1,552,769 (5.92 MB)

Non-trainable params: 512 (2.00 KB)

Optimizer params: 3,105,540 (11.85 MB)

10000/10000	<hr/>				186s 19ms/step
	precision	recall	f1-score	support	
0	0.80	0.78	0.79	159494	
1	0.79	0.81	0.80	160506	
accuracy			0.79	320000	
macro avg	0.79	0.79	0.79	320000	
weighted avg	0.79	0.79	0.79	320000	

Accuracy: 0.794375

#Implentation of Bi-LSTM and Bi-GRU model for the sentiment 140 dataset

```
from tensorflow.keras.layers import BatchNormalization, Bidirectional
model = Sequential()
model.add(Embedding(input_dim=20000, output_dim=128, input_length=50))
model.add(Bidirectional(LSTM(128, return_sequences=True)))
model.add(Dropout(0.2))
model.add(Bidirectional(GRU(128)))
model.add(Dropout(0.2))
model.add(Dense(1, activation='sigmoid'))
```

```
#Compiling the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(
```

```
#Training the model
history_4 = model.fit(X_train, y_train, epochs=5,
validation_data=(X_test, y_test), batch_size=32)
```

```
model.summary()
```

```
#Evaluate the model
predictions_5 = (model.predict(X_test) > 0.5).astype("int32")
print(classification_report(y_test, predictions_5))
print('Accuracy:', accuracy_score(y_test, predictions_5))
```

```
Epoch 1/5
40000/40000 _____ 4405s 110ms/step - accuracy: 0.7692 -
loss: 0.4763 - val_accuracy: 0.7942 - val_loss: 0.4390
Epoch 2/5
40000/40000 _____ 4131s 103ms/step - accuracy: 0.8020 -
loss: 0.4251 - val_accuracy: 0.7990 - val_loss: 0.4300
Epoch 3/5
40000/40000 _____ 3982s 100ms/step - accuracy: 0.8120 -
loss: 0.4085 - val_accuracy: 0.8000 - val_loss: 0.4315
Epoch 4/5
40000/40000 _____ 8334s 208ms/step - accuracy: 0.8180 -
loss: 0.3979 - val_accuracy: 0.7836 - val_loss: 0.4523
Epoch 5/5
40000/40000 _____ 11454s 286ms/step - accuracy: 0.8105
- loss: 0.4105 - val_accuracy: 0.7927 - val_loss: 0.4436
```

```
Model: "sequential_4"
```

Layer (type) Param #	Output Shape
embedding_4 (Embedding) 2,560,000	(32, 50, 128)

	bidirectional_5 (Bidirectional)	(32, 50, 256)	
263,168			
	dropout_10 (Dropout)	(32, 50, 256)	
0			
	bidirectional_6 (Bidirectional)	(32, 256)	
296,448			
	dropout_11 (Dropout)	(32, 256)	
0			
	dense_4 (Dense)	(32, 1)	
257			

Total params: 9,359,621 (35.70 MB)

Trainable params: 3,119,873 (11.90 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 6,239,748 (23.80 MB)

10000/10000	241s 24ms/step			
	precision	recall	f1-score	support
0	0.79	0.80	0.79	159494
1	0.80	0.78	0.79	160506
accuracy			0.79	320000
macro avg	0.79	0.79	0.79	320000
weighted avg	0.79	0.79	0.79	320000

Accuracy: 0.792653125

#Implentation of LSTM on the Davidson Dataset

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
#Loading the Davidson Dataset
```

```
data = pd.read_csv('/Users/Andrea/Desktop/Davidson Dataset.csv')
```

```
print(data.head())
```

	Unnamed: 0	count	hate_speech	offensive_language	neither	class
0	0	3	0	0	3	2
1	1	3	0	3	0	1
2	2	3	0	3	0	1
3	3	3	0	2	1	1
4	4	6	0	6	0	1

	tweet
0	!!! RT @mayasolovely: As a woman you shouldn't...
1	!!!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	!!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	!!!!!!! RT @ShenikaRoberts: The shit you...

```
data.shape
```

```
(24783, 7)
```

```
#Class Labels 0 - hate speech 1 - offensive language 2- neither
```

```
data.head(5)
```

	Unnamed: 0	count	hate_speech	offensive_language	neither	class
0	0	3	0	0	3	2
1	1	3	0	3	0	1
2	2	3	0	3	0	1
3	3	3	0	2	1	1
4	4	6	0	6	0	1

	tweet
0	!!! RT @mayasolovely: As a woman you shouldn't...
1	!!!!!! RT @mleew17: boy dats cold...tyga dwn ba...

```
2 !!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3 !!!!!!! RT @C_G_Anderson: @viva_based she lo...
4 !!!!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...
```

```
#Selecting the tweet and class column
```

```
data = data[['tweet', 'class']]
```

```
def preprocess_text(text):
    text = text.lower()
    text = re.sub(r'^a-zA-Z0-9\s', '', text)
    text = re.sub(r'\s+', ' ', text).strip()
    return text
```

```
data['tweet'] = data['tweet'].apply(preprocess_text)
```

```
label_encoder = LabelEncoder()
data['class'] = label_encoder.fit_transform(data['class'])
```

```
#Splitting the data
```

```
X_train, X_test, y_train, y_test = train_test_split(data['tweet'],
data['class'], test_size=0.2, random_state=42)
```

```
#Tokenize the text
```

```
tokenizer = Tokenizer(num_words=20000, oov_token="<OOV>")
tokenizer.fit_on_texts(X_train)
X_train_sequences = tokenizer.texts_to_sequences(X_train)
X_test_sequences = tokenizer.texts_to_sequences(X_test)
```

```
#Pad the sequences
```

```
max_length = 100
X_train_padded = pad_sequences(X_train_sequences, maxlen=max_length,
padding='post', truncating='post')
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_length,
padding='post', truncating='post')
```

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout
```

```
#Build the model
```

```
model = Sequential()
model.add(Embedding(input_dim=20000, output_dim=128,
input_length=max_length))
model.add(LSTM(128, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(128))
model.add(Dropout(0.2))
model.add(Dense(3, activation='softmax'))
```

```
#Compiling the model
```

```
model.compile(loss='sparse_categorical_crossentropy',  
optimizer='adam', metrics=['accuracy'])
```

```
#Train the model
```

```
history = model.fit(X_train_padded, y_train, epochs=5, batch_size=64,  
validation_data=(X_test_padded, y_test))
```

Epoch 1/5

```
/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/  
embedding.py:90: UserWarning: Argument `input_length` is deprecated.  
Just remove it.  
  warnings.warn(
```

```
310/310 _____ 37s 117ms/step - accuracy: 0.7632 - loss:  
0.6863 - val_accuracy: 0.7730 - val_loss: 0.6656
```

Epoch 2/5

```
310/310 _____ 36s 117ms/step - accuracy: 0.7783 - loss:  
0.6602 - val_accuracy: 0.7730 - val_loss: 0.6680
```

Epoch 3/5

```
310/310 _____ 37s 120ms/step - accuracy: 0.7719 - loss:  
0.6694 - val_accuracy: 0.7730 - val_loss: 0.6651
```

Epoch 4/5

```
310/310 _____ 35s 114ms/step - accuracy: 0.7776 - loss:  
0.6611 - val_accuracy: 0.7730 - val_loss: 0.6656
```

Epoch 5/5

```
310/310 _____ 36s 116ms/step - accuracy: 0.7726 - loss:  
0.6697 - val_accuracy: 0.7730 - val_loss: 0.6665
```

```
model.summary()
```

Model: "sequential_9"

Layer (type) Param #	Output Shape	
embedding_9 (Embedding) 2,560,000	(None, 100, 128)	
lstm_11 (LSTM) 131,584	(None, 100, 128)	
dropout_19 (Dropout) 0	(None, 100, 128)	

lstm_12 (LSTM)	(None, 128)	
131,584		
dropout_20 (Dropout)	(None, 128)	
0		
dense_9 (Dense)	(None, 3)	
387		

Total params: 8,470,667 (32.31 MB)

Trainable params: 2,823,555 (10.77 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 5,647,112 (21.54 MB)

#Evaluating the model on the test set

loss, accuracy = model.evaluate(X_test_padded, y_test)

print(f'Test Accuracy: {accuracy*100:.2f}%')

155/155 ————— 3s 22ms/step - accuracy: 0.7743 - loss: 0.6673

Test Accuracy: 77.30%

#Implementating GRU model on the Davidson twitter dataset

```
import numpy as np
import pandas as pd
import re
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, GRU, Dense, Dropout
```

#Loading the dataset

data = pd.read_csv('/Users/Andrea/Desktop/Davidson Dataset.csv')

#Using only the tweet and the class column

data = data[['tweet', 'class']]

```

#Performing preprocessing
def preprocess_text(text):
    text = text.lower()
    text = re.sub(r'^a-zA-Z0-9\s]', '', text)
    text = re.sub(r'\s+', ' ', text).strip()
    return text

data['tweet'] = data['tweet'].apply(preprocess_text)

label_encoder = LabelEncoder()
data['class'] = label_encoder.fit_transform(data['class'])

#Splitting the dataset
X_train, X_test, y_train, y_test = train_test_split(data['tweet'],
data['class'], test_size=0.2, random_state=42)

#Tokenize the text
tokenizer = Tokenizer(num_words=20000, oov_token="<OOV>")
tokenizer.fit_on_texts(X_train)
X_train_sequences = tokenizer.texts_to_sequences(X_train)
X_test_sequences = tokenizer.texts_to_sequences(X_test)

#Padding the sequences
max_length = 100
X_train_padded = pad_sequences(X_train_sequences, maxlen=max_length,
padding='post', truncating='post')
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_length,
padding='post', truncating='post')

#Building the GRU model
model = Sequential()
model.add(Embedding(input_dim=20000, output_dim=128,
input_length=max_length))
model.add(GRU(128, return_sequences=True))
model.add(Dropout(0.2))
model.add(GRU(128))
model.add(Dense(3, activation='softmax'))

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(

#Compiling the model
model.compile(loss='sparse_categorical_crossentropy',
optimizer='adam', metrics=['accuracy'])

#Training the model

```



```
history = model.fit(X_train_padded, y_train, epochs=5, batch_size=64,
validation_data=(X_test_padded, y_test))
```

Epoch 1/5

310/310 _____ 40s 125ms/step - accuracy: 0.7735 - loss: 0.6877 - val_accuracy: 0.7730 - val_loss: 0.6652

Epoch 2/5

310/310 _____ 41s 132ms/step - accuracy: 0.7807 - loss: 0.6530 - val_accuracy: 0.7730 - val_loss: 0.6658

Epoch 3/5

310/310 _____ 44s 144ms/step - accuracy: 0.7691 - loss: 0.6759 - val_accuracy: 0.7730 - val_loss: 0.6678

Epoch 4/5

310/310 _____ 42s 135ms/step - accuracy: 0.7754 - loss: 0.6647 - val_accuracy: 0.7730 - val_loss: 0.6657

Epoch 5/5

310/310 _____ 43s 140ms/step - accuracy: 0.7747 - loss: 0.6617 - val_accuracy: 0.7730 - val_loss: 0.6698

```
model.summary()
```

Model: "sequential_6"

Layer (type) Param #	Output Shape	
embedding_6 (Embedding) 2,560,000	(None, 100, 128)	
gru_5 (GRU) 99,072	(None, 100, 128)	
dropout_14 (Dropout) 0	(None, 100, 128)	
gru_6 (GRU) 99,072	(None, 128)	
dense_6 (Dense) 387	(None, 3)	

Total params: 8,275,595 (31.57 MB)

Trainable params: 2,758,531 (10.52 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 5,517,064 (21.05 MB)

#Evaluating the model

```
loss, accuracy = model.evaluate(X_test_padded, y_test, verbose=1)
print(f'Test Accuracy: {accuracy}')
```

155/155 ————— 3s 22ms/step - accuracy: 0.7743 - loss: 0.6706

Test Accuracy: 0.7730482220649719

#The below code shows the step by step procedure to implement Bi-directional LSTM and Bi-directional GRU for the Davidson twitter dataset

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

#Loading the dataset

```
df = pd.read_csv('/Users/Andrea/Desktop/Davidson Dataset.csv')
```

#Selecting the tweet and class columns

```
df= df[['tweet', 'class']]
```

```
df.head(5)
```

	tweet	class
0	!!! RT @mayasolovely: As a woman you shouldn't...	2
1	!!!!!! RT @mleew17: boy dats cold...tyga dwn ba...	1
2	!!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...	1
3	!!!!!!! RT @C_G_Anderson: @viva_based she lo...	1
4	!!!!!!! RT @ShenikaRoberts: The shit you...	1

```
print(data.head())
```

	tweet	class
0	rt mayasolovely as a woman you shouldnt compla...	2
1	rt mleew17 boy dats coldtyga dwn bad for cuffi...	1
2	rt urkindofbrand dawg rt 80sbaby4life you ever...	1
3	rt cganderson vivabased she look like a tranny	1
4	rt shenikaroberts the shit you hear about me m...	1

```

label_encoder = LabelEncoder()
data['class'] = label_encoder.fit_transform(data['class'])

#Splitting the data
X_train, X_test, y_train, y_test = train_test_split(df['tweet'],
df['class'], test_size=0.2, random_state=42)

#Tokenize the text
tokenizer = Tokenizer(num_words=20000, oov_token="<OOV>")
tokenizer.fit_on_texts(X_train)
X_train_sequences = tokenizer.texts_to_sequences(X_train)
X_test_sequences = tokenizer.texts_to_sequences(X_test)

#Pad the sequences
max_length = 100
X_train_padded = pad_sequences(X_train_sequences, maxlen=max_length,
padding='post', truncating='post')
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_length,
padding='post', truncating='post')

#Using the Bi-LSTM and Bi-GRU model for the analysis of twitter hate
speech recognition.

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Bidirectional, LSTM,
GRU, Dense, Dropout

#Defining the Bi-LSTM and Bi-GRU model
model = Sequential()
model.add(Embedding(input_dim=20000, output_dim=128,
input_length=max_length))
model.add(Bidirectional(LSTM(128, return_sequences=True)))
model.add(Dropout(0.2))
model.add(Bidirectional(GRU(128)))
model.add(Dropout(0.2))
model.add(Dense(3, activation='softmax'))

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(

#Compiling the model
model.compile(loss='sparse_categorical_crossentropy',
optimizer='adam', metrics=['accuracy'])

#Training the model
history = model.fit(X_train_padded, y_train, epochs=5, batch_size=64,
validation_data=(X_test_padded, y_test))

```

```

Epoch 1/5
310/310 _____ 76s 241ms/step - accuracy: 0.8231 - loss:
0.5061 - val_accuracy: 0.9024 - val_loss: 0.2858
Epoch 2/5
310/310 _____ 87s 281ms/step - accuracy: 0.9276 - loss:
0.2166 - val_accuracy: 0.8947 - val_loss: 0.3138
Epoch 3/5
310/310 _____ 87s 280ms/step - accuracy: 0.9573 - loss:
0.1237 - val_accuracy: 0.8933 - val_loss: 0.3579
Epoch 4/5
310/310 _____ 86s 278ms/step - accuracy: 0.9754 - loss:
0.0717 - val_accuracy: 0.8804 - val_loss: 0.5300
Epoch 5/5
310/310 _____ 87s 280ms/step - accuracy: 0.9851 - loss:
0.0473 - val_accuracy: 0.8759 - val_loss: 0.5109

```

```
model.summary()
```

```
Model: "sequential_7"
```

Layer (type) Param #	Output Shape	
embedding_7 (Embedding) 2,560,000	(None, 100, 128)	
bidirectional_7 (Bidirectional) 263,168	(None, 100, 256)	
dropout_15 (Dropout) 0	(None, 100, 256)	
bidirectional_8 (Bidirectional) 296,448	(None, 256)	
dropout_16 (Dropout) 0	(None, 256)	
dense_7 (Dense) 771	(None, 3)	

Total params: 9,361,163 (35.71 MB)

Trainable params: 3,120,387 (11.90 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 6,240,776 (23.81 MB)

#Evaluating the model on test set

```
loss, accuracy = model.evaluate(X_test_padded, y_test)
```

```
print(f'Test Accuracy: {accuracy*100:.2f}%')
```

155/155 ————— 8s 52ms/step - accuracy: 0.8785 - loss: 0.5171

Test Accuracy: 87.59%

```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['accuracy'], label='train accuracy')
```


```
plt.plot(history.history['val_accuracy'], label='val accuracy')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Accuracy')
```

```
plt.legend()
```

```
plt.show()
```



```

#This part of the code below has batch size = 32 and length if the tweet = 25

label_encoder = LabelEncoder()
data['class'] = label_encoder.fit_transform(data['class'])

#Splitting the data
X_train, X_test, y_train, y_test = train_test_split(df['tweet'],
df['class'], test_size=0.2, random_state=42)

#Tokenize the text
tokenizer = Tokenizer(num_words=20000, oov_token="<OOV>")
tokenizer.fit_on_texts(X_train)
X_train_sequences = tokenizer.texts_to_sequences(X_train)
X_test_sequences = tokenizer.texts_to_sequences(X_test)

#Pad the sequences
max_length = 25
X_train_padded = pad_sequences(X_train_sequences, maxlen=max_length,
padding='post', truncating='post')
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_length,
padding='post', truncating='post')

#Using the Bi-LSTM and Bi-GRU model for the analysis of twitter hate speech recognition.

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Bidirectional, LSTM,
GRU, Dense, Dropout

#Define the Bi-LSTM and Bi-GRU model
model = Sequential()
model.add(Embedding(input_dim=20000, output_dim=128,
input_length=max_length))
model.add(Bidirectional(LSTM(128, return_sequences=True)))
model.add(Dropout(0.2))
model.add(Bidirectional(GRU(128)))
model.add(Dropout(0.2))
model.add(Dense(3, activation='softmax'))

/opt/anaconda3/lib/python3.11/site-packages/keras/src/layers/core/
embedding.py:90: UserWarning: Argument `input_length` is deprecated.
Just remove it.
  warnings.warn(

#Compiling the model
model.compile(loss='sparse_categorical_crossentropy',
optimizer='adam', metrics=['accuracy'])

```

#Training the model

```
history = model.fit(X_train_padded, y_train, epochs=5, batch_size=32,  
validation_data=(X_test_padded, y_test))
```

Epoch 1/5

620/620 — 33s 51ms/step - accuracy: 0.8379 - loss: 0.4772 - val_accuracy: 0.9032 - val_loss: 0.2810

Epoch 2/5

620/620 — 36s 57ms/step - accuracy: 0.9334 - loss: 0.1982 - val_accuracy: 0.8927 - val_loss: 0.2988

Epoch 3/5

620/620 — 36s 58ms/step - accuracy: 0.9636 - loss: 0.1115 - val_accuracy: 0.8723 - val_loss: 0.4364

Epoch 4/5

620/620 — 36s 59ms/step - accuracy: 0.9777 - loss: 0.0671 - val_accuracy: 0.8848 - val_loss: 0.4227

Epoch 5/5

620/620 — 35s 57ms/step - accuracy: 0.9850 - loss: 0.0453 - val_accuracy: 0.8769 - val_loss: 0.5011

```
model.summary()
```

Model: "sequential_8"

Layer (type) Param #	Output Shape	
embedding_8 (Embedding) 2,560,000	(None, 25, 128)	
bidirectional_9 (Bidirectional) 263,168	(None, 25, 256)	
dropout_17 (Dropout) 0	(None, 25, 256)	
bidirectional_10 296,448	(None, 256)	
(Bidirectional)		
dropout_18 (Dropout)	(None, 256)	

0				
		dense_8 (Dense)	(None, 3)	
771				

Total params: 9,361,163 (35.71 MB)

Trainable params: 3,120,387 (11.90 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 6,240,776 (23.81 MB)

#Evaluating the model on the test set

loss, accuracy = model.evaluate(X_test_padded, y_test)

print(f'Test Accuracy: {accuracy*100:.2f}%')

155/155 ————— 2s 12ms/step - accuracy: 0.8784 - loss: 0.5116

Test Accuracy: 87.69%

import matplotlib.pyplot as plt

plt.plot(history.history['accuracy'], label='train accuracy')

plt.plot(history.history['val_accuracy'], label='val accuracy')

plt.xlabel('Epoch')

plt.ylabel('Accuracy')

plt.legend()

plt.show()

